

# AMS-512 Capital Markets and Portfolio Theory

## Fitting a Factor Model by MLE and Neural Networks Using Simulated Data.

Robert J. Frey, Research Professor  
Stony Brook University, Applied Mathematics and Statistics

Robert.Frey@StonyBrook.edu  
<http://www.ams.sunysb.edu/~frey>

---

## Overview

This notebook is a comparison of factor model fits using MLE and neural networks. The simulated data are 2,000 observations of 20 assets whose returns are driven by three factors. Half of the data is used for model fitting and half for out-of-sample testing. The factor returns and error returns in the generative model are both multivariate Normal. Both the factors and the errors are all mutually uncorrelated.

Given that there is an infinite number of ways a given space can be spanned by factors, we focus on how well the model represents the out-of-sample returns. We do, however, also discuss the differences between the two fits and make some preliminary suggestions as to how they may be reconciled.

---

## Set Up

```
In[ ]:= << Local`SaveRecall`  
      << Local`FactorFitMLE`
```

---

## Generative Model

### Mathematics

$$r_i(t) = \alpha_i + \sum_{j=1}^m b_{ij} f_j(t) + \epsilon(t) \quad i = 1, \dots, n$$

$$\mathbf{r}(t) = \boldsymbol{\alpha} + \mathbf{B}^T \mathbf{f}(t) + \boldsymbol{\epsilon}(t)$$

$$\mathbf{R} = \mathbf{A} + \mathbf{B} \mathbf{F}^T + \mathbf{E}$$

$$\mu = \alpha + \mathbf{B} \phi$$


$$\Sigma = \mathbf{B} \mathbf{C} \mathbf{B}^T + \mathbf{D}$$

## Simulation

```
In[ ]:= mnFactorReturns = RandomVariate[MultinormalDistribution[
      {0.008, 0.012, 0.010}, DiagonalMatrix[{0.2, 0.3, 0.25}^2]], 2000];
Dimensions@mnFactorReturns
```

```
Out[ ]:= {2000, 3}
```

```
In[ ]:= ListLinePlot[Accumulate@#] & /@ Transpose[mnFactorReturns]
```



```
In[ ]:= mnFactorLoadings = RandomVariate[UniformDistribution[{-0.25, 0.25}], {20, 3}];
Dimensions@mnFactorLoadings
```

```
Out[ ]:= {20, 3}
```

```
In[ ]:= mnErrors = RandomVariate[NormalDistribution[0, 0.03], {2000, 20}];
Dimensions@mnErrors
```

```
Out[ ]:= {2000, 20}
```

```
In[ ]:= vnAlphas = RandomVariate[UniformDistribution[{-0.005, 0.005}], 20];
```

```
In[ ]:= mnAssetReturns =
      vnAlphas + # & /@ (mnFactorReturns.Transpose[mnFactorLoadings] + mnErrors);
Dimensions@mnAssetReturns
```

```
Out[ ]:= {2000, 20}
```

## Training-Test Separation

```
In[ ]:= mnTraining = mnAssetReturns[[ ;; 1000]];
Dimensions@mnTraining
mnTest = mnAssetReturns[[1001 ;;]];
Dimensions@mnTest
```

```
Out[ ]:= {1000, 20}
```

```
Out[ ]:= {1000, 20}
```

## MLE Factor Model


```
In[ ]:= vnMeanMLE = Mean[mnTraining]
```

```
Out[ ]:= {-0.0091148, 0.00677535, 0.00196628, -0.00737638, -0.00350378,  
          0.00197924, 0.00338466, -0.00441579, 0.00625267, 0.00319213,  
          -0.00317594, 0.00287809, 0.00918268, -0.00860375, 0.0136834,  
          0.00512195, -0.000152026, -0.00955773, -0.00368295, 0.00459397}
```

```
In[ ]:= mnCovMLE = Covariance[mnTraining];
```

```
In[ ]:= mnCorMLE = Correlation[mnTraining];
```

```
In[ ]:= ListPlot3D[mnCorMLE, PlotRange -> All]
```




```
In[ ]:= nObs = Length[mnTraining]
```

```
Out[ ]:= 1000
```

```
In[ ]:= iOrder = 3;
```

```
In[ ]:= {{mnFactorLoadingsMLE, mnErrorMatrixMLE}, vnLLHistory, nBIC} =  
        xFactorFitMLE[nObs, mnCovMLE, xInitializeFactorModel[mnCovMLE, iOrder]];
```

```
In[ ]:= ListLinePlot[vnLLHistory, PlotRange -> All]
```



## Reconstructions

```
In[ ]:= xInverse[mnB_, mnD_] := Module[
  {mnID},
  mnID = DiagonalMatrix[1/Tr[mnD, List]];
  mnID = mnID.mnB.
  Inverse[IdentityMatrix[Last[Dimensions[mnB]]] + mnBT.mnID.mnB].mnBT.mnID
];
```

```
In[ ]:= vnFactorMeanMLE = Inverse[
  Transpose[mnFactorLoadingsMLE].Inverse[mnErrorMatrixMLE].mnFactorLoadingsMLE].
  Transpose[mnFactorLoadingsMLE].Inverse[mnErrorMatrixMLE].vnMeanMLE
```

```
Out[ ]:= {0.13442, 0.0231493, -0.0727659}
```

```
In[ ]:= mnFactorReturnsMLE =
  ((mnFactorLoadingsMLET.xInverse[mnFactorLoadingsMLE, mnErrorMatrixMLE]).
  (# - vnMeanMLE & /@ mnTraining)T) + vnFactorMeanMLE)T;
```

```
In[ ]:= vnPredMeanMLE = mnFactorLoadingsMLE.vnFactorMeanMLE
```

```
Out[ ]:= {-0.0100823, 0.0101782, 0.000535079, -0.00745652, -0.000283632,
  0.000418952, 0.00721491, -0.00552162, 0.00530291, 0.00357713,
  -0.00237757, 0.00359208, 0.00837568, -0.00590503, 0.00906559,
  0.00621896, -0.000492863, -0.00665956, -0.00548896, 0.00196353}
```


```
In[ ]:= vnAlphasMLE = vnMeanMLE - vnPredMeanMLE
```

```
Out[ ]:= {0.000967465, -0.00340283, 0.0014312, 0.0000801391, -0.00322014,
  0.00156029, -0.00383025, 0.00110583, 0.00094976, -0.000384993,
  -0.000798375, -0.000713985, 0.000807, -0.00269872, 0.00461785,
  -0.00109701, 0.000340837, -0.00289817, 0.00180601, 0.00263044}
```

# Autoencoder NN


## Topology

```
In[ ]:= netModel = NetGraph[
  {LinearLayer[3, "Input" → 20], LinearLayer[20], MeanSquaredLossLayer[]},
  {1 → 2 → NetPort["Output"], 2 → NetPort[3, "Input"],
   NetPort["Input"] → NetPort[3, "Target"], 1 → NetPort["Factor"]}
]
```

```
Out[ ]:= NetGraph[
   Input port:      vector (size: 20)
  Number of outputs: 3
  Number of layers: 3
]
```

## Initialization


```
In[ ]:= netFactor = NetInitialize[netModel]
```

```
Out[ ]:= NetGraph[
   Input port:      vector (size: 20)
  Number of outputs: 3
  Number of layers: 3
]
```

```
In[ ]:= x = RandomVariate[NormalDistribution[], {20}]
```

```
Out[ ]:= {0.521437, 0.67249, -2.04173, -0.642919, -0.916029, -0.327242,
  -0.554383, -1.27178, -1.68671, -2.46877, 2.28635, 0.108559, 0.597468,
  0.168467, -0.240104, 0.18907, 0.0244668, 0.111205, 0.0430716, 0.870715}
```

```
In[ ]:= ListPlot[{x, netFactor[x]["Output"]}]^T, PlotRange → All]
```



```
In[ ]:= netFactor[x]["Factor"]
```

```
Out[ ]:= {-1.99053, 0.642039, 1.15202}
```

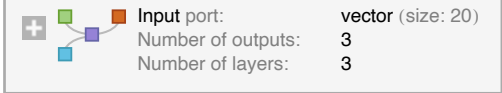
```
In[ ]:= netFactor[x][ "Loss"]
```

```
Out[ ]:= 1.59568
```

## Training

```
In[ ]:= netFactor = NetTrain[netFactor, <|"Input" → mnTraining|>]
```

```
Out[ ]:= NetGraph[
```



Input port: vector (size: 20)  
Number of outputs: 3  
Number of layers: 3

## Reconstruction


```
In[ ]:= dsInSample = Dataset[netFactor /@mnTraining];
```

```
In[ ]:= mnInSample = Normal[dsInSample[All, "Output"]];  
vnInSampleLoss = Normal[dsInSample[All, "Loss"]];  
mnInSampleFactors = Normal[dsInSample[All, "Factor"]];
```

## Comparisons (In Sample)

```
In[ ]:= mnFitMLE =  
Table[vnAlphasMLE[[i]] + mnFactorLoadingsMLE[[i]].mnFactorReturnsMLE^T, {i, 1, 20}]^T;
```

```
In[ ]:= Histogram[vnInSampleLoss, Automatic, "PDF"]
```



## Factor Correlations

Note that, as expected, the MLE produces uncorrelated factors while the NN does not.

```
In[ ]:= MatrixForm@Correlation[mnInSampleFactors]
```

```
Out[ ]//MatrixForm=
```

$$\begin{pmatrix} 1. & -0.595648 & -0.463073 \\ -0.595648 & 1. & 0.690383 \\ -0.463073 & 0.690383 & 1. \end{pmatrix}$$



```
In[ ]:= MatrixForm@Correlation[mnFactorReturnsMLE]
Out[ ]//MatrixForm=
```



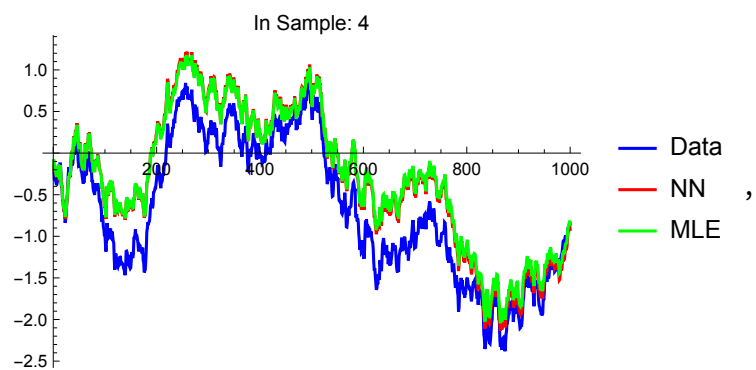
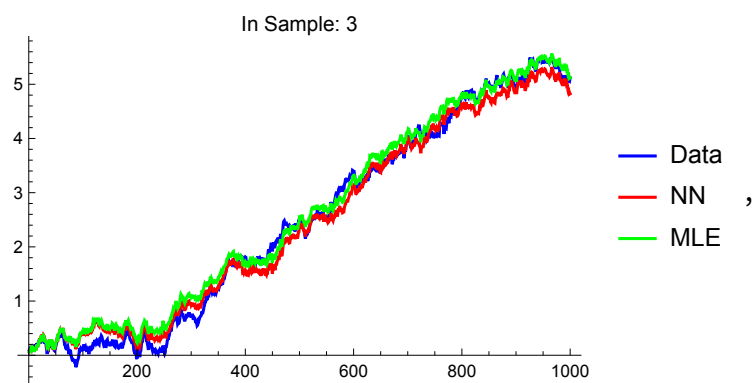
$$\begin{pmatrix} 1. & 0.000232739 & -0.000247916 \\ 0.000232739 & 1. & -0.000714157 \\ -0.000247916 & -0.000714157 & 1. \end{pmatrix}$$

## Plots





The two models produce essentially the same results when comparing the models' fits to the original data.


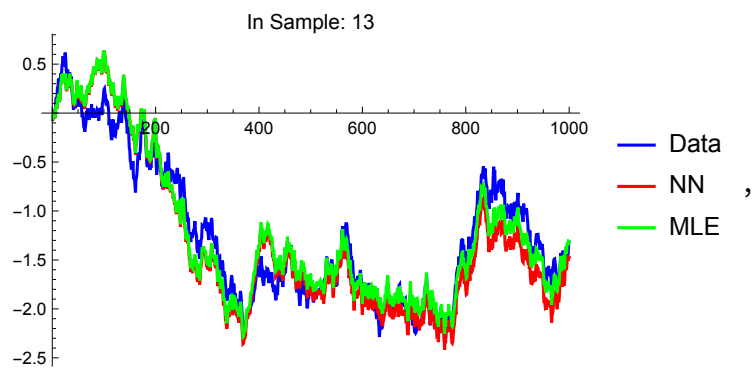

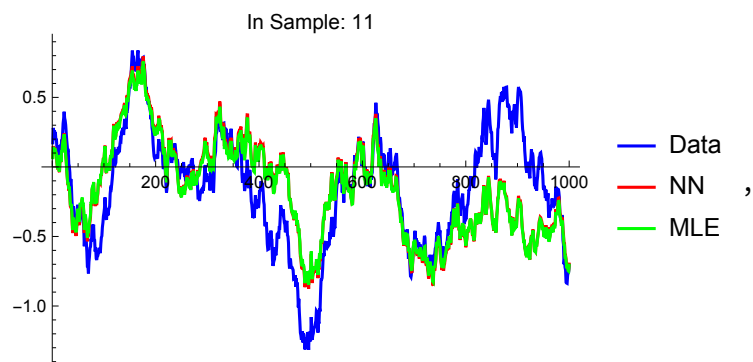
```
In[ ]:= Table[
  ListLinePlot[
    {Accumulate@mnTraining[All, i], Accumulate@mnInSample[All, i],
     Accumulate@mnFitMLE[All, i]}, PlotRange → All,
    PlotLabel → "In Sample: " <> ToString[i], PlotStyle → {{Blue}, {Red}, {Green}},
    ImageSize → 300, PlotLegends → {"Data", "NN", "MLE"}],
  {i, 1, 20}]
```




















## Out of Sample

### Factor Model MLE

```
In[ ]:= vn00SMeanMLE = Mean[mnTest]
```

```
Out[ ]:= {-0.000876323, 0.00100963, 0.00158499, 0.00281455, -0.00431954,
  0.00402954, -0.00177881, -0.00257824, 0.000249026, 0.00831942,
  -0.00309122, -0.00872581, 0.00237611, 0.00266113, 0.0000938734,
  -0.00040603, -0.00720275, -0.00315773, -0.000754216, 0.00642138}
```

```
In[ ]:= mn00SFactorReturnsMLE =
```

```
((mnFactorLoadingsMLET.xInverse[mnFactorLoadingsMLE, mnErrorMatrixMLE]).
  (# - vnMeanMLE & /@ mnTest)T) + vnFactorMeanMLE)T;
```

```
In[ ]:= vn00SPredMeanMLE = mnFactorLoadingsMLE.vnFactorMeanMLE
```

```
Out[ ]:= {0.000957328, -0.00226101, 0.00301583, -0.000221399, -0.000256967,
  -0.00274823, 0.00206702, 0.000546599, 0.00349874, 0.00267843,
  -0.000829141, -0.00447709, -0.0034064, -0.00021385, -0.000645246,
  0.00286982, -0.00190572, -0.00179925, -0.00177932, 0.000756929}
```

```
In[ ]:= vn00SA1phasMLE = vnMeanMLE - vn00SPredMeanMLE
```

```
Out[ ]:= { -0.000517777, 0.00680804, 0.00210239, -0.000632976, -0.00217255,
  0.000232772, -0.00153119, 0.000183639, -0.00164576, 0.00269598,
  0.000122069, 0.00105406, 0.0020899, 0.00292139, 0.000215715,
  -0.00154895, -0.00496882, -0.00556716, -0.00221856, 0.00280961 }
```

```
In[ ]:= mn00SFitMLE = Table[
  vn00SA1phasMLE[[i]] + mnFactorLoadingsMLE[[i]].mn00SFactorReturnsMLE^T, {i, 1, 20}]^T;
```

## Neural Network

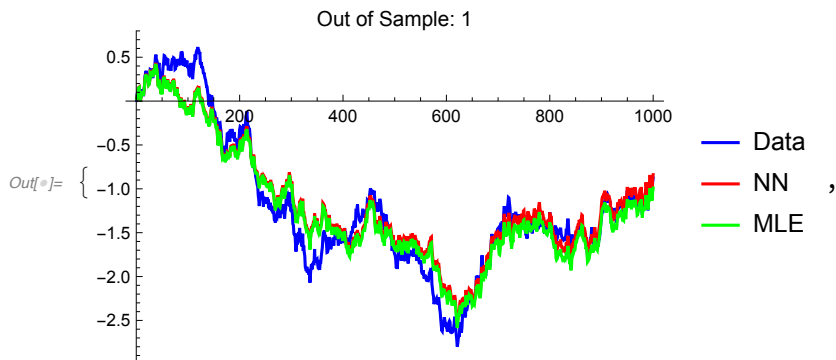
```
In[ ]:= dsOutOfSample = Dataset[netFactor /@ mnTest];
```




```
In[ ]:= mnOutOfSample = Normal[dsOutOfSample[All, "Output"]];
vnOutOfSampleLoss = Normal[dsOutOfSample[All, "Loss"]];
mnOutOfSampleFactors = Normal[dsOutOfSample[All, "Factor"]];
```




## Comparisons





### Reconstruction Plots

```
In[ ]:= Table[
  ListLinePlot[
    {Accumulate@mnTest[[All, i]], Accumulate@mnOutOfSample[[All, i]],
      Accumulate@mn00SFitMLE[[All, i]], PlotRange -> All,
      PlotLabel -> "Out of Sample: " <> ToString[i], PlotStyle -> {{Blue}, {Red}, {Green}},
      ImageSize -> 300, PlotLegends -> {"Data", "NN", "MLE"}],
  {i, 1, 20}
]
```


























## Covariances and Correlations

```
In[ ]:= mn00SCorrelations = Table[
  {Correlation[{mnTest[[All, i]], mn00SFitMLE[[All, i]]^T}][2, 1],
  Correlation[{mnTest[[All, i]], mnOutOfSample[[All, i]]^T}[2, 1]},
  {i, 1, 20}
];
```

```
In[ ]:= mn00SCovariances = Table[
  {Covariance[{mnTest[[All, i]], mn00SFitMLE[[All, i]]^T}[2, 1],
  Covariance[{mnTest[[All, i]], mnOutOfSample[[All, i]]^T}[2, 1]},
  {i, 1, 20}
];
```

```

In[ ]:= Grid[
  {{Style[Column[{"Out of Sample Correlations", "Fit vs Test"}, Center],
    FontSize → 16]},
  {TableForm[mn00SCorrelations, TableHeadings → {Automatic, {"MLE", "NN"}}]},
  Frame → All
]

```

Out[ ]:=

| Out of Sample Correlations<br>Fit vs Test |          |          |
|---|----------|----------|
|   | MLE      | NN       |
| 1   | 0.871293 | 0.870852 |
| 2   | 0.778526 | 0.776619 |
| 3   | 0.853269 | 0.850838 |
| 4   | 0.946546 | 0.946115 |
| 5   | 0.763916 | 0.764487 |
| 6   | 0.968524 | 0.968515 |
| 7   | 0.905987 | 0.906933 |
| 8   | 0.893319 | 0.890657 |
| 9   | 0.924672 | 0.925998 |
| 10  | 0.927947 | 0.922149 |
| 11  | 0.868666 | 0.870177 |
| 12  | 0.956162 | 0.959279 |
| 13  | 0.923075 | 0.922958 |
| 14  | 0.946151 | 0.947495 |
| 15  | 0.893672 | 0.895092 |
| 16  | 0.818031 | 0.82183  |
| 17  | 0.91319  | 0.915448 |
| 18  | 0.848223 | 0.846556 |
| 19  | 0.940332 | 0.939887 |
| 20  | 0.94667  | 0.947178 |

```

In[ ]:= Grid[
  {{Style[
    Column[{"Out of Sample Covariances", "Fit vs Test"}, Center], FontSize → 16}},
    {TableForm[mn00SCovariances, TableHeadings → {Automatic, {"MLE", "NN"}}]},
    Frame → All
  ]
]


```

| Out of Sample Covariances<br>Fit vs Test |            |            |
|--|------------|------------|
|  | MLE        | NN         |
| 1  | 0.00253112 | 0.0025853  |
| 2  | 0.0012271  | 0.00130834 |
| 3  | 0.00203706 | 0.00215099 |
| 4  | 0.00564936 | 0.00585918 |
| 5  | 0.00122607 | 0.00126964 |
| 6  | 0.00964287 | 0.00991296 |
| 7  | 0.00346187 | 0.00357602 |
| 8  | 0.00294983 | 0.00302496 |
| 9  | 0.00451996 | 0.00465466 |
| 10                                       | 0.0038063  | 0.0039851  |
| 11                                       | 0.00240238 | 0.00251651 |
| 12                                       | 0.00661996 | 0.0069195  |
| 13                                       | 0.00458877 | 0.00471268 |
| 14                                       | 0.00634315 | 0.0065242  |
| 15                                       | 0.00326805 | 0.00338973 |
| 16                                       | 0.00147581 | 0.00156455 |
| 17                                       | 0.00341985 | 0.00361981 |
| 18                                       | 0.00206794 | 0.00213564 |
| 19                                       | 0.00552823 | 0.00567502 |
| 20                                       | 0.00587272 | 0.00608637 |


```

In[ ]:= Show[
  ListPlot[mn00SCorrelations],
  Plot[x, {x, 0, 1}, PlotStyle → {Dashed, Red}]
]

```




```
In[ ]:= Show[
  ListPlot[mn00SCovariances],
  Plot[x, {x, 0, 1}, PlotStyle -> {Dashed, Red}]
]
```



### Comparison Mean and Standard Deviations

```
In[ ]:= Show[
  ListPlot[{Mean@mn00SFitMLE, Mean@mnOutOfSample}^T],
  Plot[x, {x, -0.12, 0.12}, PlotStyle -> {Dashed, Red}]
]
```




```
In[ ]:= LinearModelFit[{Mean@mnTest, Mean@mn00SFitMLE}^T, x, x]
```

```
Out[ ]:= FittedModel[ 0.000312706 + 0.957631 x ]
```

```
In[ ]:= LinearModelFit[{Mean@mnTest, Mean@mnOutOfSample}^T, x, x]
```

```
Out[ ]:= FittedModel[ 0.000310573 + 0.97555 x ]
```

```
In[ ]:= Show[
  ListPlot[{StandardDeviation@mn00SFitMLE, StandardDeviation@mnOutOfSample}^T,
    AxesOrigin -> {0, 0}],
  Plot[x, {x, 0, 0.1}, PlotStyle -> {Dashed, Red}, PlotRange -> All]
]
```



This first fit compares the standard deviations of the two fits.

```
In[ ]:= LinearModelFit[
  {StandardDeviation@mn00SFitMLE, StandardDeviation@mnOutOfSample}^T, x, x]
```

```
Out[ ]:= FittedModel[ 0.00120354 + 1.01688 x ]
```

Note that the fit by definition does not include the effect of errors. The next fit uses the actual data to estimate the relationship between the fit and actual standard deviation. The next two analyze how close those results are to the same computation using the out-of-sample model results.

```
In[ ]:= LinearModelFit[
  {StandardDeviation[vnAlphas + # & /@ (mnFactorReturns.Transpose[mnFactorLoadings])],
    StandardDeviation[mnAssetReturns]}^T, x, x]
```

```
Out[ ]:= FittedModel[ 0.0141635 + 0.888746 x ]
```

```
In[ ]:= LinearModelFit[
  {StandardDeviation[vnAlphas + # & /@ (mnFactorReturns.Transpose[mnFactorLoadings])],
    StandardDeviation[mnAssetReturns]}^T, x, x, IncludeConstantBasis -> False]
```

```
Out[ ]:= FittedModel[ 1.10844 x ]
```

```
In[ ]:= LinearModelFit[{StandardDeviation@mn00SFitMLE, StandardDeviation@mnTest}^T, x, x]
```

```
Out[ ]:= FittedModel[ 0.01471 + 0.897201 x ]
```

```
In[ ]:= LinearModelFit[{StandardDeviation@mn00SFitMLE, StandardDeviation@mnTest}^T,
  x, x, IncludeConstantBasis → False]
```

```
Out[ ]:= FittedModel[ 1.12723 x ]
```

```
In[ ]:= LinearModelFit[{StandardDeviation@mnOutOfSample, StandardDeviation@mnTest}^T,
  x, x, IncludeConstantBasis → False]
```

```
Out[ ]:= FittedModel[ 1.0886 x ]
```

---

## Conclusion

The bottom line is that both methods appear to be effective in modeling the returns. The advantages of the MLE approach are that its underlying theory and form are both clearly laid out and understood, including the fact that the factors it generates are mutually uncorrelated. The advantage of the neural network approach is that it is straightforward to incorporate many non-linear effects which may prove useful in improving the fit.